


Architecture — AI Screening Platform

Visl AI Labs' candidate screening platform: ingest a candidate dataset, evaluate each candidate against a job description with an LLM, analyze their GitHub at repository level, score and rank deterministically, email a test link to the shortlist, ingest test results, and schedule interviews on Google Calendar with a real Meet link — end to end, with no manual step in between.

1. High-Level Design



Why this shape. The frontend is a pure HTTP client (`frontend/lib/api_client.py` is the *only* module that talks to the backend) — it never imports backend code and holds no business logic, so the backend can be re-skinned or driven by any other client without change. The backend is stateless between requests; all state lives in Postgres, so it can be redeployed or scaled to N instances with no session affinity. Every external dependency (LLM, GitHub, email, Calendar) sits behind its own `services/` module, so a provider swap (already exercised twice — Gmail → Resend → Brevo, and Groq → Mistral on rate-limit) is a config change, not a rewrite.


2. End-to-End Workflow



Every arrow is a recruiter action through the Streamlit UI, backed by one FastAPI endpoint. Each stage persists incrementally (one candidate/row at a time, not at batch end) and degrades a single failure to a logged error rather than aborting the batch — a dead resume link or a candidate with no GitHub profile never blocks the other nine.

3. Agentic Evaluation Pipeline (LangGraph)

One compiled graph, invoked once per candidate (`backend/graph/pipeline.py`):



Node	Does	On failure
<code>load_resume</code>	Reads the candidate's extracted resume text	Missing text is not fatal — evaluation proceeds on dataset fields alone
<code>evaluate_vs_jd</code>	LLM scores skills match, project depth, experience relevance, research — each with <code>score</code> + <code>reasoning</code> + <code>evidence</code>	<code>resume_score</code> stays <code>None</code> (never 0) so weight redistributes
<code>analyze_github</code>	Fetches repos, drops forks, ranks by recency/stars/substance, sends top 3 (README + language breakdown) to the LLM alongside deterministic signals	<code>NO_PROFILE</code> / <code>NOT_FOUND</code> are normal states, not errors; only a real API failure is logged
<code>aggregate</code>	Pure Python: blends <code>resume_score</code> + <code>github_score</code> into <code>pre_test_score</code>	Never raises — computes from whatever is available

`evaluate_vs_jd` and `analyze_github` fan out in parallel and converge on `aggregate` ; every node try/excepts internally and writes to a shared `errors` list (merged via an `operator.add` reducer) instead of raising, so one candidate's exception can never take down the batch. The graph is compiled once at import and reused for all candidates.

Scoring is never LLM-computed. The LLM only emits structured per-dimension scores with reasoning; `services/scoring.py` does the arithmetic:

```
pre_test = 0.40*resume_jd + 0.25*github + 0.20*projects + 0.15*academics
final     = 0.60*pre_test  + 0.20*test_la + 0.20*test_code
```

Weights are config, not code — recruiters retune the resume/GitHub blend and re-rank instantly from stored scores, with **zero new LLM or GitHub calls**.

Resilience: Groq calls run through a round-robin pool of API keys (each an independent free-tier token bucket on the same model); a 429 rotates to the next key with a cooldown, and the pool falls back to Mistral only when every key is cooling down. `fast` mode (deployment default) favors speed by dropping to Mistral immediately; `quality` mode waits for a Groq key so every candidate is scored by the same model. Every external call (DB, GitHub, LLM, email, Calendar) is wrapped in `tenacity` retry with exponential backoff.

4. Constraints

- **Free tier only, everywhere** — Groq (rate-limited), Neon Postgres (scales to zero when idle — connections use `pool_pre_ping` + `recycle`), Render (cold starts, and outbound SMTP ports are blocked at the platform level, which is why email goes over Brevo's HTTPS API, not SMTP, in production).
- **Open-source LLM preferred** — `openai/gpt-oss-120b` on Groq is primary; Mistral is the fallback only.
- **Real Google Calendar integration, no mocks** — events must actually appear on a calendar with a working Meet link (`conferenceDataVersion=1` is mandatory; without it Google silently drops the Meet link).
- **GitHub analysis must be repository-level** — profile stats alone (followers, total stars) do not satisfy this; individual repos are fetched and their contents evaluated.
- **Dynamic schema, not a fixed CSV shape** — columns are matched case-insensitively via a synonym map; an unmappable required column fails with a named error, never a `KeyError`.
- **Python 3.11**, both apps publicly hosted, deadline-bound scope — polish is cut before Calendar integration or hosting ever is.

5. What a Recruiter Can Do

- Upload any similarly-shaped candidate CSV/XLSX — no fixed column order.
- Trigger resume download + extraction and LLM evaluation with one click each; watch live progress and partial results as a batch runs.
- Choose `fast` vs `quality` evaluation mode as a speed/consistency trade-off.
- See **every** score's reasoning and quoted evidence, not just a number — per resume dimension and per GitHub dimension.
- Re-weight the resume/GitHub blend and re-rank instantly, with **zero** inference cost, to explore "what if projects mattered more."
- Preview a shortlist by top-N or score threshold before anything is sent; send test-invite emails, with a force-resend option.
- Upload test results at any time — missing rows are never scored zero, and the final blend weights are separately tunable.
- Schedule interviews by top-N or final-score threshold, with custom working hours, slot length, gap, and timezone; cancel and re-book.
- See a per-batch pipeline-progress bar and a live dashboard (candidates, GitHub coverage, emails sent, interviews scheduled) on every page, not just the page just visited.

6. What Sets This Platform Apart

- **Explainable by construction, not by add-on** — every dimension the LLM touches carries `score` + `reasoning` + `evidence`; the UI surfaces all of it, because it was never thrown away in the first place.
- **Scoring is deterministic and auditable** — the LLM never outputs a final number. A recruiter can see exactly why two candidates are 2 points apart, and change the formula's weights without re-running inference.
- **GitHub analysis reads code, not vanity metrics** — top repos are selected by recency/stars/substance, forks dropped, README + languages sent to the LLM; a candidate with 10,000 followers and no original repos scores on substance, not reputation.

- **Provider-agnostic and already proven so** — the LLM, email, and (in design) the DB session boundary are all swappable via one config value; this was exercised for real mid-project (Gmail → Brevo, Groq → Mistral), not just designed for hypothetically.
- **Real infrastructure, not a demo stub** — Calendar events carry working Meet links today, on the free tier, with no mocked response anywhere in the path.
- **Failure is a data point, not a crash** — a dead resume link, a missing GitHub profile, a missing test score, or a throttled LLM key each degrade to an explicit state (`NO_RESULT` , `NO_PROFILE` , a logged error) that the UI shows plainly, while the rest of the batch keeps moving.